

A CREATOR TOOL THAT DOES THE WORK

Keep the context. Check the source.

From an English recording to captioned clips and a source-bound editing handoff.

Local speech recognition

Explicit review

Real video exports

Joseph Ayanda · Solo builder · September 2026

Working open-source prototype, not a subscription service.

THE MISSING-WORD DEMONSTRATION

The battery
does **not** last
twelve hours.

One omitted word changes the claim.
Compare captions against the actual audio.

SHORTER IS NOT AUTOMATICALLY FAITHFUL

Three errors can survive a polished edit.

01

The caption is wrong.

A negation, number or qualification disappears from the supplied transcript.

Listen and compare.

02

The context is outside.

The excerpt is real, but a related correction sits elsewhere in the source recording.

Inspect beyond the cut.

03

The file has changed.

An editor receives a same-name video that is not the exact media used for the edit plan.

Verify the handoff.

Initial user hypothesis: solo creators and small editing teams making source-faithful interview, tutorial and product-explanation clips.

EVIDENCE DESK

Ask the audio, not another prompt.


Whisper-tiny.en processes the actual selected audio in a browser worker. The supplied caption is not passed to the model as its answer.

01 / Generate timestamped captions from a recording, or import SRT, VTT or JSON.

02 / Compare supplied captions with recognized speech. Highlight disagreements.

03 / Review the source, edit explicitly, and keep the decision visible.

Optional model downloads require consent. Media is not uploaded. ASR can be wrong; agreement is not certification.



Actual browser inference on the fictional missing-word fixture.

FROM REVIEW TO DELIVERABLE

The edit plan belongs with the right video.

REVIEW

Select a contiguous source range.
Inspect related context.
Approve explicitly.

**BIND**

Compute media SHA-256
and byte count.
Capture the editing state.

**DELIVER**

Export captioned WebM or an edit
ZIP with an H.264/AAC renderer.

SOURCE LOCK / SAME FILE AT HANDOFF

The edit ZIP and manifest now bind to the attached media bytes. Replacing a file with the same name is not a match. SHA-256 identifies bytes, not whether the captions tell the truth.

Fingerprint attached source

Cancel

Check a previously exported manifest against the attached file **Choose File** No file chosen

MISMATCH: Different source bytes. Do not use this file with this edit manifest. Expected
6cf7c85d13ee9688658b0cf1f6c06a7bd84a89d73af79d2dc9b2326c26e1e5b3; observed
0e204ffedaec4f8e4b8d3c285552c1e0d576339e4786398466833f2caa4ee680.

Changing state invalidates old work.

Editing during hashing aborts the export. Replacing media resets reviews. Restoring matching bytes never restores approval.

The native renderer checks the source lock before output creation. A digest identifies bytes, not truth, authorship or consent.

Actual same-name, metadata-changed video rejection.

A BROWSER TOOL, NOT A REMOTE INFERENCE WRAPPER

AI proposes. The workflow preserves control.

INPUT

Local recording + captions

Web Audio decoding · SRT/VTT/JSON import · editable timed cues

ANALYZE

Speech worker

Whisper-tiny.en · pinned model revision
Transformers.js 3.8.1 + ONNX Runtime Web

Source retrieval

TF-IDF/diversity clip ranking
BM25 + lexical framing cues for context

DECIDE

Explicit review and stale-state checks

Changed transcript, selection or media invalidates affected approvals and checks.

OUTPUT

Browser rendering

Captioned portrait WebM with source audio

Source-bound edit ZIP

Timed captions · review report · manifest
Python + FFmpeg H.264/AAC route

REPRODUCIBLE INTERNAL CHECKS

A working system, with the misses visible.

153

JavaScript tests
125 inherited + 28 new

13

Native source-identity
and renderer checks

32

Browser workflows
also passed on public HTTPS

The public checks run the real app.

Actual Whisper inference, caption editing, reviewed-only export, source mismatch rejection and decoded video output. Published assets match the verified release.

The boundary detector still misses cases.

10 / 15 risky cuts flagged; **5 missed**.

2 / 9 safe cases falsely flagged.

Small internal diagnostic, not representative accuracy. Source Lock does not improve this separate score.

MAKE REVIEW EASIER, NOT INVISIBLE

Validate usefulness before claiming impact.

NOW**A complete local workflow**

Caption preparation, source-context inspection, explicit review and actual exports in one tool.

English input: up to 10 minutes / 120 MB.
Browser encoding is real-time, not frame-exact.

NEXT**Authorized creator pilots**

Compare manual editing with CutProof on consented recordings. Measure review time, missed corrections, false alarms and usable exports.

Proposed test, not completed user research.

THEN**Improve the bottleneck**

Better word alignment and retrieval; verify desktop-editor imports; test longer recordings only after memory and latency checks.

Commercial hypothesis: paid setup or team workflow support. No revenue claimed.

The value proposition is a shorter path from recording to a reviewable handoff. **Measured time savings remain unverified.**

TRY THE FAILURE CASES, NOT JUST THE HAPPY PATH

One recording. Visible evidence. A handoff you can recheck.

01 / WORKING APPLICATION

Open CutProof 1.3

02 / RECORDED DEMONSTRATION

Watch the 3:13 walkthrough

03 / CODE AND REPRODUCTION

Read source, setup and test evidence

04 / PORTABLE SOURCE

Get the tested release ZIP

Joseph Ayanda · Solo builder. Original application created September 6–8, 2026 with substantial AI assistance. Original code: MIT; third-party licenses retained. Demonstration narration: disclosed stock Kokoro synthetic voice. This is the same CutProof project also entered in AI Content Engine, not a separately invented product. No award or cash receipt is claimed.